

ECE 408 FA19 Project Report

Team: rtx5700ti

**Members: Adit Umakanth, Advaith Ravikumar,
Murugan Narayanan**

NetID: adityau2, advaith2, mrn2

UIN: 674909598, 677014085, 670975133

School Affiliation: uiuc (on campus students)

Table of Contents

Table of Contents	1
MILESTONE 1 & 2 Report	3
List of kernels that collectively consume more than 90% of the program time	3
Difference Between Kernels and API Calls	4
CPU OUTPUT	4
CPU RUNTIME	4
GPU OUTPUT	4
GPU RUNTIME	4
CPU Convolution Program Run Time	5
List of Op Times	5
MILESTONE 3 Report	6
TIMING OUTPUT	6
NVPROF OUTPUT	6
NVVP Timeline:	8
MILESTONE 4 REPORT	9
4.1 Introduction	9
4.2 Baseline	9
4.2.1 Timing Output	9
4.2.2 NVPROF Output	9
4.2.3 NVVP Timeline	11
4.2.4 Relevant Code	11
4.3 Shared Memory Convolution	11
4.3.1 Timing Output	12
4.3.2 NVPROF Output	12
4.3.3 NVVP Timeline	13
4.3.4 Relevant Code	14
4.4 Weight Matrix in Constant Memory	15
4.4.1 Timing Output	15
4.4.2 NVPROF Output	15
4.4.3 NVVP Timeline	16
4.4.4 Relevant Code	16
4.5 Multiple Kernel Implementations for Different Layer Sizes (Multikernel)	17
4.5.1 Timing Output	17
4.5.2 NVPROF Output	17
4.5.3 NVVP Timeline	18
4.5.4 Relevant Code	18

	2
4.6 Putting 3 Optimizations Together	19
4.6.1 Timing Output	19
4.6.2 NVPROF Output	19
4.6.3 NVVP Timeline	21
MILESTONE 5 (FINAL) REPORT	22
5.1 Introduction	22
5.2 Kernel Fusion for Unrolling and Matrix-Multiplication	22
5.2.1 Timing Output	22
5.2.2 NVPROF Output	23
5.2.3 NVVP Timeline	24
5.2.4 Relevant Code	24
5.3 - Exploiting Parallelism in Input Images and Output Channels	26
5.3.1 Timing Output	26
5.3.2 NVPROF Output	26
5.3.3 NVVP Timeline	28
5.3.4 Relevant Code	28
5.4 Sweeping Parameters to Find Best Values	29
5.4.1 Timing Output	30
5.4.2 NVPROF Output	30
5.4.3 NVVP Timeline	31
5.4.4 Relevant Code	31
5.4.4.1 Relevant Code from Project File	31
5.4.4.2 Python Code to plot the graphs from the excel table of values (as shown in section 5.4)	31
5.5 Tuning with Restrict and Loop Unrolling (with Multi Kernel)	32
5.5.1 Timing Output	32
5.5.2 NVPROF Output	33
5.5.3 NVVP Timeline	34
5.5.4 Relevant Code	34
5.6 Final Remarks	35
5.6.1 Overview and Kernel Description	35
5.6.2 Results	35
5.6.3 Identifying Optimization Opportunity	35
5.6.4 References	35
5.6.5 Teamwork	35
5.6.6 Directory Structure	36

MILESTONE 1 & 2 Report

List of kernels that collectively consume more than 90% of the program time

GPU activities:

31.02% 34.114ms 20 1.7057ms 1.0880us 31.926ms [CUDA memcpy HtoD]

18.02% 19.822ms 1 19.822ms 19.822ms 19.822ms
volta_scudnn_128x64_relu_interior_nn_v1

17.39% 19.123ms 4 4.7806ms 4.7767ms 4.7852ms volta_gcgemm_64x32_nt

8.71% 9.5779ms 4 2.3945ms 2.0010ms 3.1155ms void fft2d_c2r_32x32<float, bool=0, bool=0, unsigned int=0, bool=0, bool=0>(float*, float2 const *, int, int, int, int, int, int, int, int, float, float, cudnn::reduced_divisor, bool, float*, float*, int2, int, int)

7.91% 8.7044ms 1 8.7044ms 8.7044ms 8.7044ms volta_sgemm_128x128_tn

6.61% 7.2648ms 2 3.6324ms 25.504us 7.2392ms void op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7, cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*, cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float, dimArray, reducedDivisorArray)

5.82% 6.3969ms 4 1.5992ms 1.2706ms 2.0309ms void fft2d_r2c_32x32<float, bool=0, unsigned int=0, bool=0>(float2*, float const *, int, int, int, int, int, int, int, int, cudnn::reduced_divisor, bool, int2, int, int)

3.96% 4.3558ms 1 4.3558ms 4.3558ms 4.3558ms void cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)

List: [CUDA memcpy HtoD], volta_scudnn_128x64_relu_interior_nn_v1, volta_gcgemm_64x32_nt, fft2d_c2r_32x32, volta_sgemm_128x128_tn, op_generic_tensor_kernel, fft2d_r2c_32x32, cudnn::detail::pooling_fw_4d_kernel

List of all CUDA API calls that collectively consume more than 90% of the program time

API calls:

44.10% 3.35356s 22 152.43ms 14.352us 1.71565s cudaStreamCreateWithFlags

30.82% 2.34363s 24 97.651ms 75.059us 2.33580s cudaMemGetInfo

21.80% 1.65752s 19 87.238ms 1.2160us 442.33ms cudaFree

List: cudaStreamCreateWithFlags, cudaMemGetInfo, cudaFree

Difference Between Kernels and API Calls

The kernel is a function compiled and intended to be executed on a computing unit different from the CPU, such as a GPU. Once the kernel is called from the host, it is executed entirely on the GPU to completion. The results of any computation done has to be copied from/to the GPU memory into/from the CPU memory.

API calls are functions called by the host to interact with the GPU. This includes things like copying and allocating memory to the GPU or getting information about the GPU. An example of the API call is cudaMalloc or cudaFree.

CPU OUTPUT

```
* Running /usr/bin/time python m1.1.py
Loading fashion-mnist data... done
Loading model... done
New Inference
EvalMetric: {'accuracy': 0.8154}
```

CPU RUNTIME

```
19.36user 7.12system 0:10.17elapsed 260%CPU (0avgtext+0avgdata 6044932maxresident)k
0inputs+2824outputs (0major+1598360minor)pagefaults 0swaps
```

GPU OUTPUT

```
* Running /usr/bin/time python m1.2.py
Loading fashion-mnist data... done
Loading model... done
New Inference
EvalMetric: {'accuracy': 0.8154}
```

GPU RUNTIME

```
4.99user 3.18system 0:04.81elapsed 169%CPU (0avgtext+0avgdata 2973968maxresident)k
0inputs+1712outputs (0major+730591minor)pagefaults 0swaps
```

CPU Convolution Program Run Time

87.36user 11.75system 1:15.18elapsed 131%CPU (0avgtext+0avgdata 6042528maxresident)k
0inputs+0outputs (0major+2305263minor)pagefaults 0swaps

List of Op Times

Op Time: 10.779369

Op Time: 59.818806

Correctness: 0.7653 Model: ece408

MILESTONE 3 Report

TIMING OUTPUT

* Running /usr/bin/time python m3.1.py 100

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.000406

Op Time: 0.003461

Correctness: 0.76 Model: ece408

5.03user 2.84system 0:07.16elapsed 109%CPU (0avgtext+0avgdata 2783908maxresident)k

0inputs+4568outputs

(0major+634601minor)pagefaults 0swaps

* Running /usr/bin/time python m3.1.py 1000

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.004505

Op Time: 0.034755

Correctness: 0.767 Model: ece408

5.07user 3.06system 0:04.62elapsed 175%CPU (0avgtext+0avgdata 2806940maxresident)k

0inputs+0outputs (0major+640458minor)pagefaults 0swaps

* Running /usr/bin/time python m3.1.py 10000

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.051133

Op Time: 0.326978

Correctness: 0.7653 Model: ece408

5.30user 2.66system 0:05.04elapsed 157%CPU (0avgtext+0avgdata 2974616maxresident)k

0inputs+0outputs (0major+729886minor)pagefaults 0swaps

NVPROF OUTPUT

* Running /usr/bin/time python m3.1.py 10000

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.051133

Op Time: 0.326978

Correctness: 0.7653 Model: ece408

5.30user 2.66system 0:05.04elapsed 157%CPU (0avgtext+0avgdata 2974616maxresident)k

0inputs+0outputs (0major+729886minor)pagefaults 0swaps

* Running nvprof python m3.1.py

Loading fashion-mnist data... done

==528== NVPROF is profiling process 528, command: python m3.1.py

Loading model... done

New Inference

Op Time: 0.048657

Op Time: 0.324925

Correctness: 0.7653 Model: ece408

==528== Profiling application: python m3.1.py

==528== Profiling result:

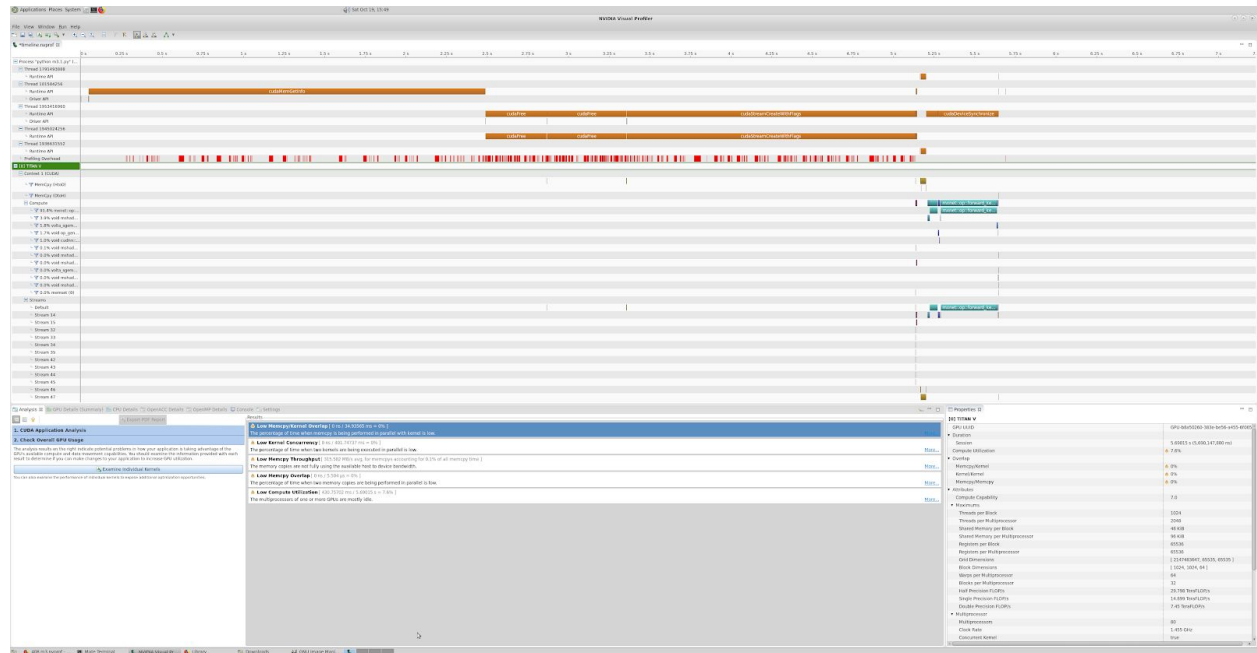
	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:								
83.39%		373.52ms	2	186.76ms	48.623ms	324.89ms		mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)
8.37%		37.510ms	20	1.8755ms	1.0560us	35.118ms		[CUDA memcpy HtoD]
3.73%		16.726ms	2	8.3632ms	3.0561ms	13.670ms		void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>, mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul, mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)								
List: forward_kernel, CUDA memcpy HostToDevice, MapPlanLargeKernel								

API calls:

40.35%	3.24361s	22	147.44ms	15.033us	1.69584s	cudaStreamCreateWithFlags
30.29%	2.43526s	22	110.69ms	72.338us	2.42973s	cudaMemGetInfo
21.12%	1.69759s	18	94.311ms	1.2300us	461.39ms	cudaFree
4.85%	390.27ms	6	65.046ms	3.2600us	324.90ms	cudaDeviceSynchronize

List: cudaStreamCreateWithFlags, cudaMemGetInfo, cudaFree, cudaDeviceSynchronize

NVVP Timeline:



The kernel we wrote only accounts for 345 ms out of the total runtime of about 5.6 s. It was interesting to see that the majority of the execution time is not our kernel, but rather `cudaStreamCreateWithFlags` and `cudaMemGetInfo`.

MILESTONE 4 REPORT

4.1 Introduction

For this milestone, we chose to tackle 3 of the optimizations in the suggested list on the project GitHub page. We will explain these optimizations and show the effect of said optimizations on the runtime of the neural network.

It is important to note that all our optimizations output the correct accuracy as mentioned on the project GitHub page, so this means that there is no change to the output of each convolution operation.

As recommended, we implemented these optimizations independently to better showcase the effect that each one has on the runtime of the convolution. The optimizations we chose to implement are as follows:

- Shared Memory Convolution
- Weight Matrix (Kernel Values) in Constant Memory
- Multiple Kernel Implementations for Different Layer Sizes

4.2 Baseline

Our baseline kernel without any optimizations is a bit different from our kernel from Milestone 3. We are now using the threads in the z dimension to process different images in parallel, and each thread is looping through the image channels instead. We were doing the opposite in Milestone 3. This change has significantly improved our execution time.

4.2.1 Timing Output

Op Time: 0.058203

Op Time: 0.140968

Correctness: 0.7653 Model: ece408

4.2.2 NVPROF Output

==353== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	73.59%	199.12ms		2	99.561ms	58.172ms	140.95ms
mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)							
	12.64%	34.199ms	20	1.7099ms	1.0560us	32.092ms	[CUDA memcpy HtoD]
	6.28%	16.983ms	2	8.4917ms	3.0878ms	13.896ms	void mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>,

```
mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul,
mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>,
float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)
```

```
2.93% 7.9166ms 1 7.9166ms 7.9166ms 7.9166ms
```

```
volta_sgemm_128x128_tn
```

```
2.71% 7.3307ms 2 3.6653ms 25.312us 7.3054ms void
```

```
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7,
cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*,
cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float,
dimArray, reducedDivisorArray)
```

```
1.63% 4.4010ms 1 4.4010ms 4.4010ms 4.4010ms void
```

```
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,
cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *,
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,
cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float,
cudnnPoolingStruct, int, cudnn::reduced_divisor, float)
```

```
API calls: 43.30% 3.19750s 22 145.34ms 29.914us 1.59818s
```

```
cudaStreamCreateWithFlags
```

```
30.67% 2.26454s 22 102.93ms 69.659us 2.25993s cudaMemGetInfo
```

```
21.54% 1.59030s 18 88.350ms 1.4450us 417.40ms cudaFree
```

```
2.93% 216.12ms 6 36.021ms 2.6080us 140.96ms
```

```
cudaDeviceSynchronize
```

```
0.94% 69.293ms 9 7.6993ms 26.442us 32.284ms cudaMemcpy2DAsync
```

```
0.27% 19.702ms 29 679.37us 2.1270us 11.120ms
```

```
cudaStreamSynchronize
```

```
0.17% 12.268ms 66 185.88us 9.0930us 4.5967ms cudaMalloc
```

```
0.07% 5.4763ms 4 1.3691ms 867.92us 1.8577ms
```

```
cudaGetDeviceProperties
```

```
0.03% 2.4806ms 375 6.6140us 384ns 334.65us
```

```
cuDeviceGetAttribute
```

```
0.02% 1.3124ms 2 656.21us 643.51us 668.91us cudaHostAlloc
```

```
0.02% 1.3044ms 912 1.4300us 468ns 19.038us
```

```
cudaFuncSetAttribute
```

4.2.3 NVVP Timeline



4.2.4 Relevant Code

Kernel code:

```
for (int m = 0; m < M; m++){    // Loop through images
    conv_val = 0.0; // Reset val for new image
    for (int c = 0; c < C; c++){    // Loop through input features
        // Loop through 2D convolution kernel/filter
        for (int kx = 0; kx < K; kx++){
            for (int ky = 0; ky < K; ky++){
                if ((tz < B) && (ty + ky < H) && (tx + kx <
W)){ // if dimensions are within input bounds and M
                    conv_val += k4d(m, c, ky, kx) * x4d(tz, c,
ty+ky, tx+kx);    // do convolution
                }
            }
        }
    }
}
```

4.3 Shared Memory Convolution

The main bottleneck with a lot of GPU algorithms is the bandwidth on global memory. Constantly writing and reading data to and from the global memory slows down the access time for each thread and can cut down the theoretical speed of algorithms by orders of magnitude.

Using shared memory is one way to get around this bottleneck when there is a high level of data reuse, like in convolution. Each thread block has its own shared memory. Data is copied here

from global memory, but is only done so once, at the start of the kernel function. Then for subsequent reads, data can be fetched from the shared memory instead of from global memory.

This is advantageous for multiple reasons. Each thread block only accesses its own shared memory, so there are no bandwidth issues unlike when all thread blocks in the kernel try to read and write to global memory. The second advantage is that shared memory is much faster than global memory, even without taking bandwidth issues into consideration. According to the [Nvidia Developer Blog](#), “Because it is on-chip, shared memory is much faster than local and global memory. In fact, shared memory latency is roughly 100x lower than uncached global memory latency “.

Similar to MP4, we used strategy 2 to load the input image snippet that was going to be used in the convolution. We looped through input channels and got the (X,Y, Image) position from the thread indices to load the correct input addresses. Similar to the MP, we load $C \cdot (\text{TILE_WIDTH} + K - 1)^2$ input elements, where C is the number of input channels, and K is the convolution filter size. Each element is used M times in our convolution as we loop through the output channels.

4.3.1 Timing Output

Op Time: 0.030706

Op Time: 0.103325

Correctness: 0.7653 Model: ece408

4.3.2 NVPROF Output

==352== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	66.36%	133.95ms	2	66.976ms	30.655ms	103.30ms	
mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)	16.25%	32.797ms	20	1.6398ms	1.0240us	30.606ms	[CUDA memcpy HtoD]
	7.34%	14.819ms	2	7.4093ms	2.9247ms	11.894ms	void mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>, mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul, mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)
volta_sgemv_128x128_tn	4.01%	8.1008ms	1	8.1008ms	8.1008ms	8.1008ms	
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7, cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*, cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float, dimArray, reducedDivisorArray)	3.57%	7.2106ms	2	3.6053ms	25.216us	7.1854ms	void

```

2.17% 4.3736ms    1 4.3736ms 4.3736ms 4.3736ms void
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,
cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *,
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,
cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float,
cudnnPoolingStruct, int, cudnn::reduced_divisor, float)


```

```

API calls: 42.98% 3.12682s    22 142.13ms 14.367us 1.60838s
cudaStreamCreateWithFlags
    30.77% 2.23869s    22 101.76ms 68.516us 2.23426s cudaMemGetInfo
21.41% 1.55756s    18 86.531ms 1.1890us 418.63ms cudaFree
    2.05% 148.79ms    6 24.799ms 2.8080us 103.30ms
cudaDeviceSynchronize
    0.91% 66.412ms    9 7.3791ms 29.860us 30.750ms cudaMemcpy2DAsync
    0.79% 57.600ms    66 872.72us 6.2700us 32.962ms cudaMalloc
    0.48% 34.815ms    912 38.174us    444ns 8.1835ms
cudaFuncSetAttribute
    0.27% 19.741ms    29 680.73us 2.2350us 10.992ms
cudaStreamSynchronize
    0.13% 9.3560ms    6 1.5593ms 1.8700us 9.3257ms cudaEventCreate
    0.07% 4.9837ms    4 1.2459ms 436.31us 1.8680ms
cudaGetDeviceProperties
    0.03% 2.4735ms    375 6.5950us    407ns 332.83us
cuDeviceGetAttribute
    0.02% 1.4020ms    216 6.4900us 1.2480us 162.43us
cudaEventCreateWithFlags

```

4.3.3 NVVP Timeline



4.3.4 Relevant Code

Kernel code:

```
__shared__ float shared_mem[MAX_C*SM_DIM*SM_DIM];
...
// copy shared memory
// Images * In_Channels * Y * X;

for (int c = 0; c < C; c++){
    int index = (c*SM_DIM*SM_DIM) + (ty*SM_DIM) + (tx);
    // if (loc_x_i==0 && loc_y_i==0){
    //     printf("IDX: %d, C: %d %d, tzyx = %d %d %d, X:%d, Y:%d Z:
%d\n", index, c, C, tz, ty, tx, loc_x_i, loc_y_i, loc_z_i);
    // }
    // __syncthreads();
    if (loc_z_i < B && loc_y_i < H && loc_x_i < W){
        shared_mem[index] = x4d(loc_z_i, c, loc_y_i, loc_x_i);
    }
    else{
        shared_mem[index] = 0.0;
    }
}

__syncthreads();

// Convolution Code
if ((ty < TILE_WIDTH) && (tx < TILE_WIDTH)){ // if dimensions
are within input bounds and M
    for (int m = 0; m < M; m++){ // Loop through output features
        float conv_val = 0.0; // Reset val for new output element
        for (int c = 0; c < C; c++){ // Loop through input
features
            // Loop through 2D convolution kernel/filter
            for (int ky = 0; ky < K; ky++){
                for (int kx = 0; kx < K; kx++){
                    int index = c*SM_DIM*SM_DIM + (ty+ky)*SM_DIM +
(tx+kx);
                    conv_val += k4d(m, c, ky, kx) *
shared_mem[index]; // do convolution
                }
            }
        }
    }
}
```

4.4 Weight Matrix in Constant Memory

Similar to shared memory, CUDA provides another interesting feature to circumvent the bandwidth restriction on access to global memory. Unlike shared memory, constant memory is not on-chip, but is cached and designed to be accessed faster than global memory, in certain situations.

Constant memory slowed down our kernel. We suspect that the overhead of copying the weights from global to constant memory was greater than the speed boost we would've gotten from reading from global memory. The weights in global memory would also be stored in a relatively fast L1 cache, thus resulting in no performance gains. It is worth noting, however, that the first layer (M=12, C = 1) had a slight speedup, but the second one did not (M=24, C=12). This is probably because the smaller layer had less cudaMemcpyToSymbol overhead.

We can clearly see from the NVVP timeline that the constant memory kernel is considerably slower than our other optimized kernels.

4.4.1 Timing Output

Op Time: 0.059142

Op Time: 0.872424

Correctness: 0.7653 Model: ece408

4.4.2 NVPROF Output

==353== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	93.22%	931.42ms	2	465.71ms	59.049ms	872.37ms	
mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)							
	3.30%	32.971ms	20	1.6485ms	1.0240us	30.755ms	[CUDA memcpy HtoD]
	1.47%	14.706ms	2	7.3532ms	2.9287ms	11.778ms	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>,							
mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul,							
mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>,							
float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)							
	0.79%	7.8598ms	1	7.8598ms	7.8598ms	7.8598ms	
volta_sgemm_128x128_tn							
	0.72%	7.2170ms	2	3.6085ms	25.376us	7.1917ms	void
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7,							
cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*,							
cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float,							
dimArray, reducedDivisorArray)							


```

0.44% 4.3831ms      1 4.3831ms 4.3831ms 4.3831ms void
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,
cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *,
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,
cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float,
cudnnPoolingStruct, int, cudnn::reduced_divisor, float)

```

```

API calls: 39.00% 3.23393s      22 147.00ms 10.292us 1.65115s
cudaStreamCreateWithFlags
27.98% 2.32000s      22 105.45ms 72.683us 2.31533s cudaMemGetInfo
19.42% 1.61049s      18 89.472ms 1.2540us 431.41ms cudaFree
11.41% 946.15ms       6 157.69ms 4.3470us 872.37ms cudaDeviceSynchronize
0.81% 66.956ms       9 7.4395ms 27.056us 30.968ms cudaMemcpy2DAsync
0.48% 39.627ms      66 600.41us 6.5380us 16.713ms cudaMalloc
0.45% 37.312ms     912 40.912us 439ns 19.195ms cudaFuncSetAttribute
0.23% 19.362ms      29 667.66us 3.2070us 10.927ms cudaStreamSynchronize

```

4.4.3 NVVP Timeline



4.4.4 Relevant Code

```

__constant__ float mask[24*12*5*5];
...
conv_val += mask[tz*C*K*K+c*K*K+ky*K+kx] * x4d(b, c, ty+ky, tx+kx);

```

Host code:

```

cudaMemcpyToSymbol(mask, k.dptr_, M*C*K*K*sizeof(float), 0,
cudaMemcpyHostToDevice);

```

4.5 Multiple Kernel Implementations for Different Layer Sizes (Multikernel)

We have two kernel functions, one for each layer size. The “small” kernel is for the M=12, C=1 layer, and the “large” kernel is for the M=24, C=12 layer. The two kernels have appropriate sizes for shared memory, and we see a performance boost for this. By allocating the right amount of resources, we waste less time writing unnecessary 0's to the shared mem or checking bounds and reduce our op time.

4.5.1 Timing Output

Op Time: 0.030934

Op Time: 0.105163

Correctness: 0.7653 Model: ece408

4.5.2 NVPROF Output

==352== Profiling result:

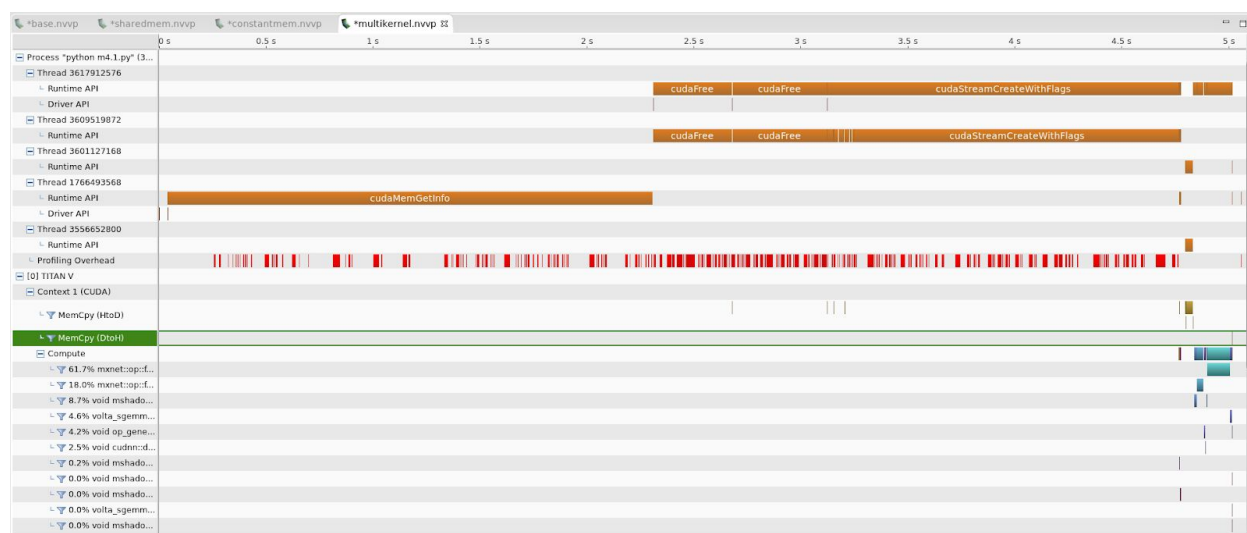
Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	51.70%	105.13ms	1	105.13ms	105.13ms	105.13ms	
mxnet::op::forward_kernel_large(float*, float const *, float const *, int, int, int, int, int, int)	15.90%	32.331ms	20	1.6166ms	1.0560us	30.302ms	[CUDA memcopy HtoD]
	15.19%	30.883ms	1	30.883ms	30.883ms	30.883ms	
mxnet::op::forward_kernel_small(float*, float const *, float const *, int, int, int, int, int, int)	7.32%	14.891ms	2	7.4455ms	2.9188ms	11.972ms	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>, mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul, mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)	3.91%	7.9609ms	1	7.9609ms	7.9609ms	7.9609ms	
volta_sgemm_128x128_tn	3.54%	7.2062ms	2	3.6031ms	24.576us	7.1816ms	void
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7, cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*, cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float, dimArray, reducedDivisorArray)	2.14%	4.3506ms	1	4.3506ms	4.3506ms	4.3506ms	void
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,							

cudaNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)

API calls:

43.27%	3.12584s	22	142.08ms	14.882us	1.59030s	cudaStreamCreateWithFlags
	31.03%	2.24181s	22	101.90ms	67.657us	2.23739s cudaMemGetInfo
	21.27%	1.53674s	18	85.375ms	1.2290us	413.20ms cudaFree
	2.09%	150.93ms	6	25.155ms	3.8290us	105.14ms
cudaDeviceSynchronize						
	0.91%	65.532ms	9	7.2813ms	28.552us	30.455ms cudaMemcpy2DAsync
	0.68%	48.780ms	912	53.487us		431ns 39.469ms
cudaFuncSetAttribute						
	0.27%	19.572ms	29	674.90us	3.0820us	10.936ms
cudaStreamSynchronize						
	0.22%	15.810ms	66	239.54us	6.4660us	3.0577ms cudaMalloc

4.5.3 NVVP Timeline



4.5.4 Relevant Code

```
#define TILE_WIDTH 16
#define MASK_WIDTH 5
#define MAX_M 24
#define MIN_M 12
#define MAX_C 12
#define MIN_C 1
#define SM_DIM (TILE_WIDTH+MASK_WIDTH-1) // shared memory dim
...
```

```

__global__ void forward_kernel_large(float *y, const float *x, const
float *k, const int B, const int M, const int C, const int H, const
int W, const int K)
{
    __shared__ float shared_mem[MAX_C*SM_DIM*SM_DIM];
...
__global__ void forward_kernel_small(float *y, const float *x, const
float *k, const int B, const int M, const int C, const int H, const
int W, const int K)
{
    __shared__ float shared_mem[MIN_C*SM_DIM*SM_DIM];

```

Host code:

```

// Call the kernel
    if (C==1){
        forward_kernel_small<<<gridDim, blockDim,
0>>>(y.dptr_,x.dptr_,k.dptr_, B,M,C,H,W,K);
    }
    else{
        forward_kernel_large<<<gridDim, blockDim,
0>>>(y.dptr_,x.dptr_,k.dptr_, B,M,C,H,W,K);
    }

```

4.6 Putting 3 Optimizations Together

We combined all of the optimizations that we have made thus far into one kernel and ran it to see our overall improvement.

4.6.1 Timing Output

Op Time: 0.027426

Op Time: 0.233445

Correctness: 0.7653 Model: ece408

4.6.2 NVPROF Output

==352== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	69.78%	233.38ms	1	233.38ms	233.38ms	233.38ms	
mxnet::op::forward_kernel_large(float*, float const *, float const *, int, int, int, int, int, int)	11.61%	38.835ms	20	1.9417ms	1.0880us	36.254ms	[CUDA memcpy HtoD]

```

      8.17% 27.324ms      1 27.324ms 27.324ms 27.324ms
mxnet::op::forward_kernel_small(float*, float const *, float const *, int, int, int, int, int, int)
      4.45% 14.878ms      2 7.4390ms 2.9155ms 11.962ms void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>,
mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul,
mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>,
float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)
      2.35% 7.8720ms      1 7.8720ms 7.8720ms 7.8720ms
volta_sgemm_128x128_tn
      2.15% 7.2061ms      2 3.6031ms 24.832us 7.1813ms void
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7,
cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*,
cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float,
dimArray, reducedDivisorArray)
      1.30% 4.3555ms      1 4.3555ms 4.3555ms 4.3555ms void
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,
cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *,
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,
cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float,
cudnnPoolingStruct, int, cudnn::reduced_divisor, float)
      0.12% 400.83us      1 400.83us 400.83us 400.83us void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int,
mshadow::Shape<int=2>, int=2, int)

```

```

API calls: 42.09% 3.40399s      22 154.73ms 15.150us 1.76411s
cudaStreamCreateWithFlags
      30.02% 2.42783s      22 110.36ms 71.383us 2.42554s cudaMemGetInfo
      21.39% 1.73013s      18 96.118ms 1.2840us 473.91ms cudaFree
      3.41% 275.61ms       6 45.934ms 3.1460us 233.39ms cudaDeviceSynchronize
      0.96% 78.031ms       9 8.6701ms 30.357us 36.332ms cudaMemcpy2DAsync
      0.80% 64.431ms      66 976.22us 5.8590us 46.459ms cudaMalloc
      0.57% 46.384ms     912 50.859us 452ns 24.523ms cudaFuncSetAttribute
      0.35% 28.082ms     216 130.01us 1.3010us 26.555ms
cudaEventCreateWithFlags
      0.24% 19.713ms      29 679.76us 3.6400us 10.895ms cudaStreamSynchronize

```


MILESTONE 5 (FINAL) REPORT

5.1 Introduction

For the milestone, we chose to implement 3 more of the optimizations in the suggested list on the project GitHub page. We will explain these optimizations and show the effect of said optimizations on the runtime of the neural network.

Once again, it is important to note that all our optimizations output the correct accuracy as mentioned on the project GitHub page, so this means that there is no change to the output of each convolution operation.

As recommended, we implemented these optimizations independently to better showcase the effect that each one has on the runtime of the convolution. Instead of just comparing our improvements to the baseline kernel mentioned in Section 4.2 of this report, we are adding optimizations on top of each other and tweaking parameters to find the true improvement. This is because parameters in a certain optimization might be different when applied to the base kernel instead of a kernel with other optimizations to yield the greatest speedup. We chose to implement the following optimizations:

- Kernel fusion for unrolling and matrix-multiplication
- Exploiting parallelism in input images and output channels
- Sweeping various parameters to find best values
- Tuning with restrict and loop unrolling

5.2 Kernel Fusion for Unrolling and Matrix-Multiplication

To increase performance, instead of having a separate kernel to unroll the filter and input matrices, we perform unrolling when loading the tile into shared memory in the matrix multiplication kernel. This is done by calculating the data indices such that the matrices are unrolled as they are loaded into shared memory. Thus, the work previously done by 2 kernels is fused into a single kernel. This improves performance because we eliminate the overhead needed to launch an extra kernel, as well as allocating and writing to extra memory to facilitate transferring the unrolled matrices between the two kernels. As shown in the NVVP timeline, we just have the overhead for launching 1 kernel instead of 2 kernels and extra memory operations.

5.2.1 Timing Output

Op Time: 0.045572

Op Time: 0.064764

Correctness: 0.7653 Model: ece408

5.2.2 NVPROF Output

==265== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	60.76%	110.22ms	2	55.110ms	45.501ms	64.720ms	
mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)							
	18.76%	34.030ms	20	1.7015ms	1.0880us	31.852ms	[CUDA memcpy HtoD]
	9.33%	16.924ms	2	8.4618ms	3.0649ms	13.859ms	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>,							
mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul,							
mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>,							
float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)							
	4.37%	7.9183ms	1	7.9183ms	7.9183ms	7.9183ms	
volta_sgemv_128x128_tn							
	4.00%	7.2644ms	2	3.6322ms	24.767us	7.2396ms	void
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7,							
cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*,							
cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float,							
dimArray, reducedDivisorArray)							
	2.43%	4.4041ms	1	4.4041ms	4.4041ms	4.4041ms	void
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,							
cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *,							
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,							
cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float,							
cudnnPoolingStruct, int, cudnn::reduced_divisor, float)							
	0.24%	430.94us	1	430.94us	430.94us	430.94us	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,							
mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int,							
mshadow::Shape<int=2>, int=2, int)							
	0.04%	68.415us	1	68.415us	68.415us	68.415us	void
mshadow::cuda::SoftmaxKernel<int=8, float,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>>(mshadow::gpu,							
int=2, unsigned int)							
	0.04%	64.863us	13	4.9890us	1.1840us	24.447us	void
mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,							
mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int,							
mshadow::Shape<int=2>, int=2)							

API calls:	42.04%	3.34171s	22	151.90ms	14.106us	1.71621s	
cudaStreamCreateWithFlags							
	32.90%	2.61493s	22	118.86ms	72.908us	2.60953s	cudaMemGetInfo
20.84%	1.65646s	18	92.026ms	1.2460us	442.58ms		cudaFree
	1.60%	127.22ms	6	21.203ms	12.163us	64.734ms	cudaDeviceSynchronize
	0.86%	68.219ms	9	7.5798ms	24.591us	31.856ms	cudaMemcpy2DAsync
	0.82%	65.263ms	216	302.14us	1.2870us	42.583ms	
cudaEventCreateWithFlags							
	0.30%	23.722ms	912	26.010us	440ns	10.917ms	cudaFuncSetAttribute
	0.25%	19.869ms	29	685.13us	5.6940us	11.026ms	cudaStreamSynchronize
	0.14%	11.345ms	66	171.90us	6.4640us	1.7428ms	cudaMalloc
	0.09%	7.4554ms	4	1.8639ms	1.2294ms	2.4802ms	
cudaGetDeviceProperties							
	0.05%	3.6243ms	12	302.02us	7.1360us	3.0850ms	cudaMemcpy
	0.03%	2.7242ms	375	7.2640us	413ns	382.66us	cuDeviceGetAttribute
	0.02%	1.2958ms	9	143.98us	9.4150us	1.0333ms	cudaMemsetAsync
	0.01%	1.0011ms	2	500.56us	54.646us	946.48us	cudaHostAlloc
	0.01%	741.00us	27	27.444us	10.200us	68.685us	cudaLaunchKernel
	0.01%	626.13us	4	156.53us	97.395us	276.55us	cuDeviceTotalMem
	0.01%	566.84us	4	141.71us	75.494us	248.02us	cudaStreamCreate
	0.01%	435.44us	202	2.1550us	823ns	44.130us	cudaDeviceGetAttribute

5.2.3 NVVP Timeline



5.2.4 Relevant Code

Kernel code:

```
...
int numMatACols = C*K*K; // = numMatBRows
...
// Convolution Code
int num_iter = ceil(numMatACols/(1.0*TILE_WIDTH));
```

```

for (int i = 0; i < num_iter; i++)
{
    int temp_col = i * TILE_WIDTH + tx;
    int temp_row = i * TILE_WIDTH + ty;

    // filter tensor
    int k_m = row;
    int k_c = temp_col / (K*K);
    int k_h = (temp_col % (K*K)) / K;
    int k_w = (temp_col % (K*K)) % K;

    // input tensor
    int x_b = b;
    int x_c = temp_row / (K*K);
    int x_p = (temp_row % (K*K)) / K;
    int x_q = (temp_row % (K*K)) % K;
    int x_h = col / W_out;
    int x_w = col % W_out;

    if (temp_col < numMatACols && row < M)
        tileMatA[ty][tx] = k4d(k_m, k_c, k_h, k_w);
    else
        tileMatA[ty][tx] = 0.0;

    if (temp_row < numMatACols && col < H_out * W_out)
        tileMatB[ty][tx] = x4d(x_b, x_c, x_h + x_p, x_w + x_q);
    else
        tileMatB[ty][tx] = 0.0;

    __syncthreads();

    for (int q = 0; q < TILE_WIDTH; q++)
        conv_val += tileMatA[ty][q] * tileMatB[q][tx];

    __syncthreads();
}

// output tensor
int y_b = b;
int y_m = row;
int y_h = col / W_out;
int y_w = col % W_out;

if (row < M && col < W_out * H_out)

```

```
y4d(y_b, y_m, y_h, y_w) = conv_val;
```

Host code:

```
...
dim3 blockDim(TILE_WIDTH, TILE_WIDTH, 1);
dim3 gridDim(ceil((H_OUT*W_OUT)/(1.0*TILE_WIDTH)),
ceil((M)/(1.0*TILE_WIDTH)), B);
// call kernel
...
```

5.3 - Exploiting Parallelism in Input Images and Output Channels

We exploited parallelism in our optimized kernel by threading on M (number of output channels), and B (number of input images). As seen in our kernel launch code, for both sizes of kernels, the Y dimension of the grid depends on a tiled parameter of M, and the Z dimension of the grid is B.

We wanted to exploit parallelism on input images since: (1) there are 10,000 images, so looping through B in the kernel would be extremely slow (what we did in milestone 3. Section 3.1 shows that the total Op Time was around 370 ms). (2) each image is the same size, so the same kind of work must be done on it, making this an ideal parameter to exploit parallelly. (3) The dimension of the input is B * C * H * W. Thus, we get no memory coalescence by having the same thread work on multiple images.

We wanted to exploit parallelism on output channels since the same work needs to be done on every output channel. We leveraged this by setting out tile width according to the M values which were known to be 12 in the first layer (for the small kernel) and 24 in the second layer (for the large kernel). This arrangement is very intuitive and worked really well. Each thread writes to the output and no threads need to be turned off. Thus it is resource efficient.

We had actually noticed this potential improvement after completing Milestone 3, so we implemented using the threads in the Z dimension to process different images in parallel, and using threads to process the image channels in parallel instead of each thread processing channels through for loops. This became our new baseline kernel as discussed in Section 4.2. In this section, we optimized this further by taking advantage of tiling.

5.3.1 Timing Output

Op Time: 0.026548

Op Time: 0.046994

Correctness: 0.7653 Model: ece408

5.3.2 NVPROF Output

==266== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	32.14%	46.975ms	1	46.975ms	46.975ms	46.975ms	
mxnet::op::forward_kernel_large(float*, float const *, float const *, int, int, int, int, int, int)	24.27%	35.469ms	20	1.7734ms	1.1200us	33.154ms	[CUDA memcpy HtoD]
	18.13%	26.499ms	1	26.499ms	26.499ms	26.499ms	
mxnet::op::forward_kernel_small(float*, float const *, float const *, int, int, int, int, int, int)	11.57%	16.909ms	2	8.4544ms	3.0651ms	13.844ms	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>,							
mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul,							
mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>,							
float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)	5.42%	7.9183ms	1	7.9183ms	7.9183ms	7.9183ms	
volta_sgemm_128x128_tn							
	5.02%	7.3307ms	2	3.6654ms	25.696us	7.3050ms	void
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7,							
cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*,							
cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float,							
dimArray, reducedDivisorArray)	3.01%	4.3994ms	1	4.3994ms	4.3994ms	4.3994ms	void
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,							
cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *,							
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,							
cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float,							
cudnnPoolingStruct, int, cudnn::reduced_divisor, float)	0.31%	452.48us	1	452.48us	452.48us	452.48us	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,							
mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int,							
mshadow::Shape<int=2>, int=2, int)							

API calls:

41.71%	3.23193s	22	146.91ms	14.966us	1.65991s	cudaStreamCreateWithFlags
32.69%	2.53254s	22	115.12ms	68.282us	2.52803s	cudaMemGetInfo
21.74%	1.68419s	18	93.566ms	1.2200us	441.91ms	cudaFree
1.17%	90.400ms	6	15.067ms	3.0110us	46.979ms	cudaDeviceSynchronize
0.92%	71.309ms	9	7.9232ms	38.307us	33.192ms	cudaMemcpy2DAsync
0.80%	62.119ms	912	68.112us	439ns	11.141ms	cudaFuncSetAttribute
0.40%	30.729ms	66	465.59us	6.6360us	17.906ms	cudaMalloc
0.26%	20.016ms	29	690.19us	2.4070us	11.102ms	cudaStreamSynchronize
0.12%	9.3476ms	12	778.96us	6.7800us	8.9002ms	cudaMemcpy
0.09%	6.8663ms	4	1.7166ms	426.36us	2.2094ms	
						cudaGetDeviceProperties

0.03%	2.6189ms	375	6.9830us	411ns	333.61us	cuDeviceGetAttribute
0.02%	1.8248ms	216	8.4470us	1.2690us	466.65us	cudaEventCreateWithFlags

5.3.3 NVVP Timeline



5.3.4 Relevant Code

```
#define TILE_WIDTH_LARGE 24
#define TILE_WIDTH_SMALL 12
```

Host code:


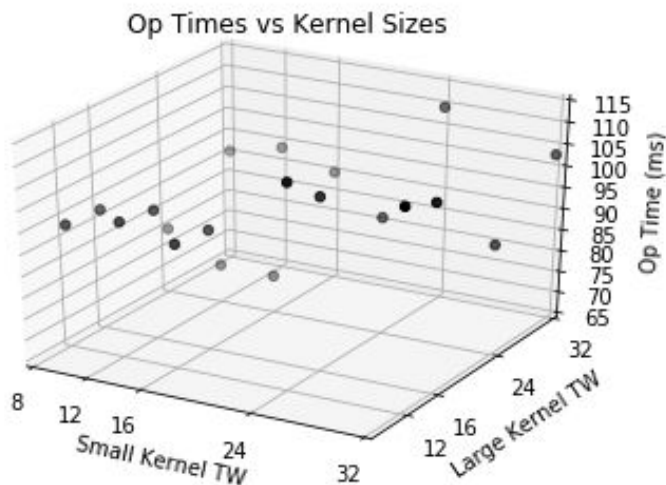
```
// Call the kernel
if (C==1){
    dim3 blockDim(TILE_WIDTH_SMALL, TILE_WIDTH_SMALL, 1);
    dim3 gridDim(ceil((H_OUT*W_OUT)/(1.0*TILE_WIDTH_SMALL)),
    ceil((M)/(1.0*TILE_WIDTH_SMALL)), B);

    // Call the kernel
    forward_kernel_small<<<gridDim, blockDim,
0>>>(y.dptr_,x.dptr_,k.dptr_, B,M,C,H,W,K);
}
else{
    dim3 blockDim(TILE_WIDTH_LARGE, TILE_WIDTH_LARGE, 1);
    dim3 gridDim(ceil((H_OUT*W_OUT)/(1.0*TILE_WIDTH_LARGE)),
    ceil((M)/(1.0*TILE_WIDTH_LARGE)), B);

    // Call the kernel
    forward_kernel_large<<<gridDim, blockDim,
0>>>(y.dptr_,x.dptr_,k.dptr_, B,M,C,H,W,K);
}
```

5.4 Sweeping Parameters to Find Best Values

To improve our runtime, we swept through parameters for the small and large tile widths for the corresponding kernel, as implemented from our multi-kernel optimization. We tried tile widths of sizes 8, 12, 16, 24, and 32 for the small kernel and sizes 12, 16, 24, 32 for the large kernel. We chose to test tile widths that were powers of 2 and a couple of values in between to maximize the usage of the tiles. As shown in the graphs below, the overall op time was lowest when using a small tile width of 16 and a large tile width of 24. We saw in section 5.3 how small tile width of 12 and large tile width of 24 would be a good choice, but found that having the small tile width as 16 improved our performance. This makes sense because the boost from increased reuse of the input matrix elements (by having more threads to load the input matrix and larger shared memory) outweighs the cost of having 4 threads that do not write to the output channels. These parameters gave us our best Op Time.



Kernel		Op Times (s)		
Small Size	Large Size	Op Small	Op Large	Total
8	12	0.029922	0.062788	0.092710
8	16	0.029910	0.062204	0.092114
8	24	0.030098	0.049235	0.079333
8	32	0.029897	0.060317	0.090214
12	12	0.026274	0.069462	0.095736
12	16	0.026441	0.067913	0.094354
12	24	0.026124	0.046834	0.072958
12	32	0.026427	0.066855	0.093282
16	12	0.023350	0.069720	0.093070
16	16	0.023492	0.068635	0.092127
16	24	0.023500	0.049244	0.072744
16	32	0.023353	0.066219	0.089572
24	12	0.042873	0.068766	0.111639
24	16	0.042870	0.061578	0.104448
24	24	0.042641	0.048913	0.091554
24	32	0.042608	0.066660	0.109268
32	12	0.040896	0.070158	0.111054
32	16	0.040681	0.067166	0.107847
32	24	0.040911	0.049059	0.089970
32	32	0.040678	0.062044	0.102722

Best Params			
16	24	0.0235	0.049244
Ranking Op Time (ms)			70.684

5.4.1 Timing Output

Op Time: 0.023880

Op Time: 0.047964

Correctness: 0.7653 Model: ece408

5.4.2 NVPROF Output

==264== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	34.42%	47.944ms	1	47.944ms	47.944ms	47.944ms	
mxnet::op::forward_kernel_large(float*, float const *, float const *, int, int, int, int, int, int)	23.00%	32.029ms	20	1.6015ms	1.0880us	29.884ms	[CUDA memcpy HtoD]
	17.12%	23.848ms	1	23.848ms	23.848ms	23.848ms	
mxnet::op::forward_kernel_small(float*, float const *, float const *, int, int, int, int, int, int)	10.61%	14.777ms	2	7.3886ms	2.9253ms	11.852ms	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>, mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul, mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)	6.10%	8.4950ms	1	8.4950ms	8.4950ms	8.4950ms	
volta_sgemm_128x128_tn	5.17%	7.2056ms	2	3.6028ms	25.376us	7.1802ms	void
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7, cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*, cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float, dimArray, reducedDivisorArray)	3.13%	4.3573ms	1	4.3573ms	4.3573ms	4.3573ms	void
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)	0.29%	409.50us	1	409.50us	409.50us	409.50us	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2, int)							

API calls:

41.63%	3.05899s	22	139.05ms	15.055us	1.56945s	cudaStreamCreateWithFlags
33.60%	2.46880s	22	112.22ms	70.062us	2.46418s	cudaMemGetInfo

20.99%	1.54239s	18	85.688ms	1.2250us	408.67ms	cudaFree
1.18%	86.585ms	6	14.431ms	2.8350us	47.948ms	cudaDeviceSynchronize
0.88%	64.902ms	9	7.2113ms	28.238us	30.037ms	cudaMemcpy2DAsync
0.74%	54.707ms	912	59.985us	431ns	26.032ms	cudaFuncSetAttribute
0.29%	21.393ms	66	324.13us	6.6920us	4.8030ms	cudaMalloc
0.27%	20.134ms	29	694.28us	2.4010us	10.955ms	cudaStreamSynchronize
0.25%	18.046ms	216	83.548us	1.2650us	8.0995ms	
cudaEventCreateWithFlags						
0.06%	4.2093ms	4	1.0523ms	339.42us	1.5823ms	
cudaGetDeviceProperties						
0.03%	2.1682ms	375	5.7810us	410ns	273.44us	cuDeviceGetAttribute

5.4.3 NVVP Timeline



5.4.4 Relevant Code

5.4.4.1 Relevant Code from Project File

```
#define TILE_WIDTH_LARGE 24
#define TILE_WIDTH_SMALL 16
```

5.4.4.2 Python Code to plot the graphs from the excel table of values (as shown in section 5.4)

```
# imports
from mpl_toolkits.mplot3d import axes3d
import matplotlib.pyplot as plt
import numpy as np

# data array copied over from the excel sheet
data = np.array([
    [[ 8],[12],[0.09271 ]],[[ 8],[16],[0.092114]],[[ 8],[24],[0.079333]],
    [[ 8],[32],[0.090214]],[[12],[12],[0.095736]],[[12],[16],[0.094354]],
    [[12],[24],[0.072958]],[[12],[32],[0.093282]],[[16],[12],[0.09307 ]],
    [[16],[16],[0.092127]],[[16],[24],[0.072744]],[[16],[32],[0.089572]],
```



```

[[24],[12],[0.111639]], [[24],[16],[0.104448]], [[24],[24],[0.091554]],
[[24],[32],[0.109268]], [[32],[12],[0.111054]], [[32],[16],[0.107847]],
[[32],[24],[0.08997 ]], [[32],[32],[0.102722]]])

X = data[:,0]          # tile widths for small kernel
Y = data[:,1]          # tile widths for large kernel
Z = data[:,2]*1000     # total op times in ms

# plot graph
fig = plt.figure()
ax = fig.add_subplot(111,projection='3d')
ax.scatter3D(X,Y,Z, c='black', marker='o')
# label graph and save it
ax.set_xlabel('Small Kernel TW')
ax.set_ylabel('Large Kernel TW')
ax.set_zlabel('Op Time (ms)')
ax.set_xlim3d([8,32])
ax.set_ylim3d([8, 32])
ax.set_zlim3d([65,115])
ax.set_xticks([8,12,16,24,32])
ax.set_yticks([12,16,24,32])
ax.set_zticks([65,70,75,80,85,90,95,100,105,110,115])
ax.set_title('Op Times vs Kernel Sizes')
plt.savefig('graph_3d.png')

```

5.5 Tuning with Restrict and Loop Unrolling (with Multi Kernel)

We used this in the convolutional kernel to great effect. By unrolling the inner-most loop, we were able to get a better runtime (it performed better than multi-kernel in section 4.5). The inner loop was for the X dim of the convolution filter. The size is always K=5, so we copied the multiplication and addition to `conv_val` statement 5 times. We found that manually copying the code performed better than `#pragma unroll` by a few milliseconds (timing not provided as it was only slightly faster). We think that this is because the programmer explicitly stating what they want out of the unrolling is better than the compiler guessing what the programmer wants in this case. It was also a small loop. We tried adding `#pragma unroll` to other loops individually and concurrently, but the initially described configuration performed the best. This optimization was not very helpful for the fusion matrix multiply kernel. This is probably because shared memory matrix multiply is already a fairly efficient algorithm and doesn't get much gain from loop unrolling since there aren't many loops (direct `matmul` vs looping to multiply).

5.5.1 Timing Output

New Inference

Op Time: 0.014937

Op Time: 0.070421

5.5.2 NVPROF Output

==264== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	45.51%	70.393ms	1	70.393ms	70.393ms	70.393ms	
mxnet::op::forward_kernel_large(float*, float const *, float const *, int, int, int, int, int, int)	21.79%	33.702ms	20	1.6851ms	1.0880us	31.490ms	[CUDA memcpy HtoD]
	9.62%	14.879ms	1	14.879ms	14.879ms	14.879ms	
mxnet::op::forward_kernel_small(float*, float const *, float const *, int, int, int, int, int, int)	9.58%	14.822ms	2	7.4110ms	2.9220ms	11.900ms	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>, mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul, mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)	5.63%	8.7039ms	1	8.7039ms	8.7039ms	8.7039ms	
volta_sgemm_128x128_tn	4.66%	7.2078ms	2	3.6039ms	25.823us	7.1820ms	void
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7, cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*, cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float, dimArray, reducedDivisorArray)	2.82%	4.3576ms	1	4.3576ms	4.3576ms	4.3576ms	void
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)	0.26%	406.59us	1	406.59us	406.59us	406.59us	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2, int)							
API calls:	41.20%	3.27691s	22	148.95ms	15.324us	1.70117s	
cudaStreamCreateWithFlags	33.69%	2.67948s	22	121.79ms	71.907us	2.67536s	cudaMemGetInfo
	20.80%	1.65423s	18	91.902ms	1.2490us	442.30ms	cudaFree
	1.35%	107.36ms	216	497.06us	1.3130us	105.32ms	
cudaEventCreateWithFlags	1.26%	100.13ms	6	16.688ms	4.2680us	70.400ms	cudaDeviceSynchronize
	0.86%	68.378ms	9	7.5976ms	30.821us	31.717ms	cudaMemcpy2DAsync

0.25%	20.082ms	29	692.50us	3.4460us	10.709ms	cudaStreamSynchronize
0.23%	18.202ms	912	19.958us	440ns	9.2931ms	cudaFuncSetAttribute
0.19%	15.319ms	66	232.11us	6.6730us	4.8858ms	cudaMalloc
0.07%	5.2087ms	4	1.3022ms	434.08us	1.9509ms	
cudaGetDeviceProperties						
0.03%	2.6270ms	375	7.0050us	413ns	339.11us	cuDeviceGetAttribute

5.5.3 NVVP Timeline



5.5.4 Relevant Code

Kernel code:

```
__global__ void forward_kernel_small(float* __restrict__ y, const
float* __restrict__ x, const float* __restrict__ k, const int B,
const int M, const int C, const int H, const int W, const int K)
{
    ...
    for (int ky = 0; ky < K; ky++){
        conv_val += k4d(m, c, ky, 0) * shared_mem[index];
        // do convolution
        conv_val += k4d(m, c, ky, 1) *
        shared_mem[index+1]; // do convolution
        conv_val += k4d(m, c, ky, 2) *
        shared_mem[index+2]; // do convolution
        conv_val += k4d(m, c, ky, 3) *
        shared_mem[index+3]; // do convolution
        conv_val += k4d(m, c, ky, 4) *
        shared_mem[index+4]; // do convolution
        index += SM_DIM_SMALL;
    }
    ...
}
```

5.6 Final Remarks

5.6.1 Overview and Kernel Description

Overall, by implementing various optimizations, we were able to greatly increase the performance of our image convolution kernel. Our best kernel was the one which combined Kernel Fusion for Unrolling and Matrix-Multiplication (section 5.2) with shared memory (incorporated in section 5.2 and also used in section 4.3), multiple kernels for different layer sizes (section 4.5), and incorporating values found from sweeping (section 5.4). Section 5.4 builds off of the “exploiting parallelism” optimization described in section 5.3. This kernel is also described above in some detail in section 5.4. We called this the “Multi-fusion” kernel.

5.6.2 Results

Of the 7 optimizations we implemented, using constant memory and restrict plus loop unrolling slowed down the total op time of the multi-fusion kernel, so we left them out of our final version used to compete for speed in the rankings. Our unoptimized kernel from Milestone 3 had op times of $0.051133 + 0.326978 = 0.378111$ seconds. Our most optimal kernel, Multi-fusion (described above), had op times of $0.023880 + 0.047964 = 0.071844$ seconds, while maintaining full correctness. That is approximately a 5.263x speedup.

5.6.3 Identifying Optimization Opportunity

We mainly referred to the list of possible optimizations given on the project GitHub page. We related them to course material and labs from the semester and set about implementing them. All relevant details are provided in the writeup provided for each optimization. We discussed our approach to them, the results, and justifications.

5.6.4 References

We used course material (lecture slides, past exams, textbook) for this project. If we used any external sources, they have been mentioned.

5.6.5 Teamwork

The team functioned smoothly. The initial checkpoints were very straight-forward so the team met up and collaborated to finish them (code, report, nvprof, nvvp). In the later checkpoints, the team members tried different optimizations and we picked whichever ones were successful. We organized the team into code and report. Advait focussed more on the code, and Adit and Murugan focussed more on the report and nvvp, as this had more components. The team would be flexible and help write/debug code or finish/submit the report as necessary. This was true for all of the optimizations.

5.6.6 Directory Structure

We have included all of our optimizations in the submission file under the ece408_src folder.

The optimizations can be found in the following files:

- 1) Shared Memory convolution: shared_mem.cuh, new-forward.cuh, multi-kernel.cuh, restrict_loop_unroll.cuh
- 2) Kernel fusion for unrolling and matrix-multiplication: kernel_fusion.cuh, new-forward.cuh
- 3) Weight matrix (kernel values) in constant memory: constant_mem.cuh
- 4) Tuning with restrict and loop unrolling: restrict_loop_unroll.cuh
- 5) Sweeping various parameters to find best values: new-forward.cuh
- 6) Exploiting parallelism in input images, input channels, and output channels: new-forward.cuh, multi_kernel.cuh
- 7) Multiple kernel implementations for different layer sizes: new-forward.cuh, multi_kernel.cuh

Our baseline kernel is included as base.cuh